
MORPH-DA: A Mutation-Grounded Benchmark for Metamorphic Verification of Data-Analysis Agents

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Data-analysis agents increasingly generate and execute Python programs, but
2 successful execution does not establish that the program implemented the re-
3 requested filters, aggregation, grouping, denominator, or date window. We introduce
4 MORPH-DA, a benchmark and environment-grounded runtime verifier for such
5 *wrong-but-executable* programs. The verifier is gold-answer-free at runtime but
6 specification-conditioned: it receives a structured analytical specification, never
7 the reference output, and checks operator-aware metamorphic relations by rerun-
8 ning the same program on controlled data transformations. The benchmark contains
9 101 compositional tasks and 563 validated, non-equivalent semantic mutants.
10 MORPH-DA detects 364/563 faults (64.7%, 95% task-clustered CI [58.4, 70.5])
11 versus 9/563 for universal robustness checks ($p = 9.4 \times 10^{-79}$). Re-executing
12 each seed-42 agent program unchanged on held-out data seeds reveals that 13–
13 16 programs per model (20–26% of apparently correct programs) are *accidental*
14 *corrects*. With this corrected ground truth, MORPH-DA obtains 79–89% precision
15 and 67–79% recall on executable programs, with 8–21% false-positive rate.
16 A one-step repair study on 91 wrong programs yields 12.1% fixes from relation-
17 specific feedback versus 5.5% from generic feedback. These results position meta-
18 morphic execution as a complementary, zero-additional-LLM-call verifier rather
19 than a correctness certificate.

20 1 Introduction

21 LLM-powered data-analysis agents translate natural-language questions into executable programs
22 over tabular data. An execution monitor catches syntax errors, missing columns, and runtime excep-
23 tions, but a program can run cleanly while performing the wrong analysis: omitting a segment filter,
24 using a sum instead of a mean, counting rows instead of distinct entities, grouping by the wrong
25 dimension, or swapping current and comparison periods. We call these *wrong-but-executable pro-*
26 *grams* (WEPs). Their outputs are especially dangerous because they are plausible, reproducible, and
27 accompanied by no runtime signal.

28 Gold-answer evaluation identifies WEPs in a benchmark, but the correct output is rarely available
29 during deployment and gives little diagnostic information for repair. A second problem is less
30 visible: a structurally wrong program can coincidentally return the correct answer on one finite
31 dataset. For example, omitting a cancelled-order filter may not change the top category in one data
32 realization but can change it after the category mix shifts. We call such cases *accidental corrects*.
33 They make single-instance accuracy an optimistic estimate of whether an agent implemented the
34 requested computation.

35 MORPH-DA addresses these problems through metamorphic runtime verification. Instead of asking
36 whether the original answer matches an unavailable gold output, it asks whether the same generated

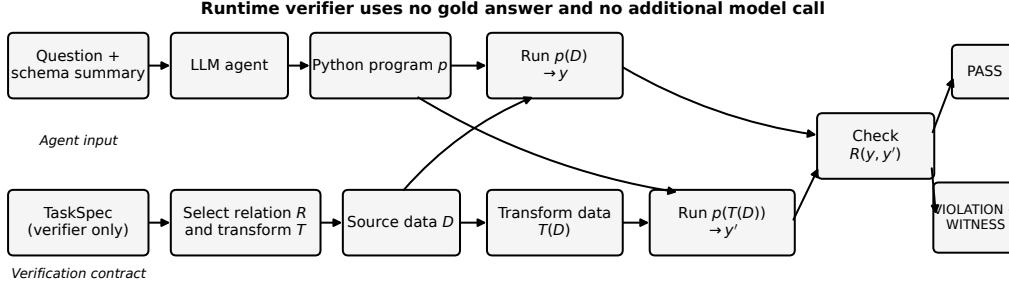


Figure 1: MORPH-DA separates program generation from verification. The agent receives only the question and schema summary. The verifier receives the candidate program and a structured task specification, generates deterministic follow-up tables, re-executes the unchanged program, and checks the expected output relation. No gold answer or additional LLM call is used during verification.

program behaves predictably after controlled changes to the data. Adding extreme rows outside a requested date range should not change the answer; duplicating all rows should double a sum but preserve a mean; perturbing a non-median extreme should change a mean but not a median; and constructing a dominant counterfactual group should force the expected winner. A violated relation produces a concrete counterexample witness.

Our contributions are:

1. **A mutation-grounded verification benchmark.** We provide 101 compositional analytical tasks and 563 validated non-equivalent mutants across aggregation, filtering, grouping, ranking, and hardcoding faults.
2. **Operator-aware, environment-grounded verification.** More than 20 relations transform the underlying data and check invariance, equivariance, monotonicity, or a forced counterfactual outcome without exposing the gold answer.
3. **An accidental-correct evaluation protocol.** We hold each generated program fixed and re-execute it on held-out data distributions, revealing a substantial gap between single-seed and structural correctness.
4. **Controlled and natural-agent evidence.** We report mutation coverage, natural-program precision/recall and accepted-answer risk, multi-transformation sensitivity, and a paired one-step repair study.

2 Related Work

Data-analysis agents and benchmarks. InfiAgent-DABench evaluates end-to-end analysis over CSV files [Hu et al., 2024]. DS-1000 and DSCoDeBench test data-science code generation [Lai et al., 2022, Ouyang et al., 2025]. Recent benchmarks broaden coverage to agentic workflows, heterogeneous workspaces, and end-to-end computer environments [Zhang et al., 2025, Nam et al., 2026, Li et al., 2026, Sun et al., 2026, Rahman et al., 2026]. These works primarily score final task success. MORPH-DA instead evaluates whether an execution-grounded verifier detects semantic faults, including programs that happen to return the correct output on one dataset.

Metamorphic and mutation testing. Metamorphic testing validates necessary relationships across transformed inputs when a direct test oracle is unavailable [Segura et al., 2016]. Mutation testing injects controlled faults to quantify a test suite’s sensitivity [Jia and Harman, 2011, Petrović et al., 2021]. Prior work applies metamorphic transformations to text-to-SQL systems and language models [Ma and Wang, 2020, Yang et al., 2025, Cho et al., 2025]. MORPH-DA differs in transforming the *tabular environment* of an unchanged agent-generated Python program, conditioning relations on analytical operators, and using a validated semantic-mutant corpus to measure coverage.

71 **Verification signals for agents.** Model-based judges and self-critique can inspect a trace, but they
 72 introduce an additional model call and may share the generator’s semantic blind spots. MORPH-DA
 73 supplies a deterministic, falsifiable signal: a specific transformed dataset and pair of outputs that
 74 violate a declared relation. It is complementary to model-based judging rather than a replacement
 75 for it.

76 3 Problem Setting

77 A task is $t = (q, D, s)$, where q is a natural-language question, D is a dictionary of DataFrames, and
 78 s is a structured analytical specification. The agent sees q and a schema summary, and produces a
 79 reusable program p implementing `analyze(tables)`. The verifier sees p and s , but not a reference
 80 program or gold output.

81 The source result is $y = p(D)$. A metamorphic test $m = (T, R)$ contains a deterministic data
 82 transformation T and a necessary output relation R . It produces $D' = T(D)$ and $y' = p(D')$. A
 83 violation is

$$v_m(p, t) = \mathbf{1}[\neg R_m(y, y')]. \quad (1)$$

84 The program is flagged if any applicable relation violates. Relations express equality, known scaling,
 85 monotonic change, or a forced winner. Passing all relations is not proof of correctness: each relation
 86 encodes only a necessary condition.

87 For controlled evaluation, a semantic mutant is *killed* when at least one relation flags it. Mutation
 88 score is the fraction of validated, non-equivalent mutants killed. For natural programs, we report
 89 precision, recall, false-positive rate (FPR), and accepted-answer risk $\text{AAR} = \text{FN}/(\text{FN} + \text{TN})$, the
 90 probability that an accepted executable program is wrong.

91 4 MORPH-DA

92 4.1 Specification-conditioned relation library

93 Each TaskSpec records filters, metric operators, grouping dimensions, date windows, ranking di-
 94 rection, comparison periods, and output type. The natural-language question is generated from this
 95 specification, but the agent never receives the specification itself. This separation lets the benchmark
 96 test whether an agent inferred and implemented the requested computation, while the verifier checks
 97 compliance without knowing the numeric answer.

98 Relations are activated only when their premises are justified by the specification. Representative
 99 families are:

- 100 • **Universal robustness:** row permutation and column-order invariance detect accidental po-
 101 sitional dependence.
- 102 • **Filter and scope:** inject extreme rows that violate a required filter; a correct program must
 103 ignore them. A complementary in-scope sentinel can be constructed to force a winner.
- 104 • **Aggregation and statistics:** full duplication preserves means, medians, distinct counts,
 105 ratios, and rankings but scales sums and counts. Non-median extreme perturbations dis-
 106 tinguish mean from median. Translation, scaling, and count-vs-distinct checks provide
 107 additional operator-specific tests.
- 108 • **Grouping and ranking:** counterfactual group construction, rank reversal, and top- k con-
 109 sistency expose wrong grouping dimensions, ascending/descending errors, and hardcoded
 110 labels.
- 111 • **Time and period comparisons:** period-isolated injections and a forced year-over-year
 112 winner expose swapped periods, missing date windows, and absolute-vs-relative change
 113 errors.

114 The current library supports sum, mean, median, count, distinct count, min, max, standard deviation,
 115 variance, quantile, ratio, percentage change, and correlation. Fields for weighted means and joins
 116 exist in the specification DSL, but no weighted-mean or join relations are implemented or evaluated
 117 in the present corpus.

118 4.2 Counterexample witnesses and repair

119 When a relation fails, the verifier records the transformation, source output, follow-up output, ex-
120 pected relation, and likely fault family. This record is a concrete counterexample rather than an
121 opaque confidence score. In the repair experiment, the same base program is paired across four
122 feedback strategies: no retry, generic feedback, the violated relation name, and a full witness.

123 4.3 Multi-transformation sensitivity

124 Many transformations contain randomized but deterministic choices, such as which eligible rows
125 become sentinels. Running a second verifier random seed generates a different set of follow-up
126 tables while keeping the candidate program and source data fixed. We use this as a sensitivity
127 analysis, not as a uniformly preferred deployment mode, because additional tests can increase both
128 recall and false positives.

129 5 Benchmark Construction

130 **Tasks and data.** The benchmark contains 101 single-table tasks over eight synthetic business sce-
131 narios: retail orders, web sessions, seller marketplace, SaaS subscriptions, marketing campaigns,
132 payments, fulfillment, and customer support. Tasks span five compositional levels: scalar aggrega-
133 tion (24), grouped ranking (48), ratios or support thresholds (11), year-over-year comparison (16),
134 and multi-filter ratio-plus-comparison tasks (2). Data generators include heavy-tailed metrics, nulls,
135 duplicates, and current/prior periods. The same schema and question are retained while seeds change
136 the underlying distribution.

137 **Filter discriminability.** For each required filter, we verify that applying it changes the gold answer
138 on seeds 7, 42, and 123. Three tasks fail this discriminability check and are excluded from the 98-
139 task cross-seed evaluation set; otherwise, a program omitting the filter could pass regardless of
140 implementation.

141 **Reference compiler and consistency.** A structural compiler translates each TaskSpec into a Pan-
142 das reference program. All references execute on seeds 7, 42, and 123 and pass applicable relations.
143 Because references and relations are co-developed from the same specification, this establishes
144 benchmark internal consistency, not an independent zero-FPR estimate. Independent false-positive
145 behavior is measured on LLM-generated programs.

146 **Mutation funnel.** AST operators generate 3,030 candidate mutations. Applicability, syntax, ex-
147 ecution, and output-contract filters leave 709 candidates for equivalence testing. Across five oracle
148 seeds, 146/709 (20.6%) are equivalent and removed, yielding 563 validated non-equivalent mutants:
149 131 aggregation, 173 filter, 21 grouping, 129 hardcoding, and 109 ranking faults.

150 6 Experiments

151 **Agents.** A LangGraph code agent generates one Python program per task from the question and
152 schema summary using Claude Haiku 4.5, Sonnet 4.6, or Opus 4.5. Evaluation centers on the 101
153 seed-42 programs per model. The framework is an implementation substrate, not part of the claimed
154 method.

155 **Three evaluation conditions.** Condition A uses seed-42 gold labels on all 101 tasks and therefore
156 treats accidental corrects as correct. Condition B saves each seed-42 program and re-executes that
157 exact source unchanged on seeds 7 and 123; failure on either reclassifies it as wrong. Conditions B
158 and C use the 98 filter-discriminating tasks. Condition C retains the same corrected ground truth but
159 runs MORPH-DA with two transformation seeds instead of one. Precision, recall, and FPR exclude
160 original execution failures; their counts are reported separately.

161 **Statistics.** We use continuity-corrected McNemar tests for paired detector comparisons, Holm
162 correction across the three natural-agent comparisons, and a task-clustered bootstrap with 3,000

Table 1: Detection on 563 validated, non-equivalent mutants. The intermediate configuration is a genuine run recorded alongside the full result.

Verifier	Killed	Score	95% CI
Execution only	0	0.0%	–
Universal relations	9	1.6%	–
Universal + filter + aggregation	346	61.5%	–
Full MORPH-DA	364	64.7%	[58.4, 70.5]

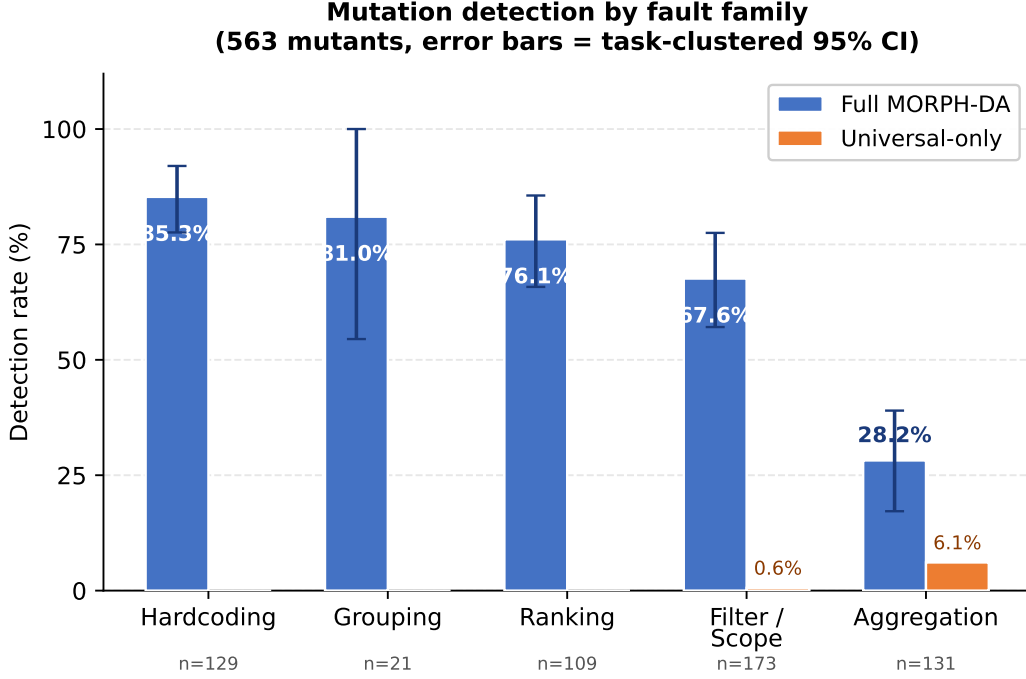


Figure 2: Mutation detection by fault family. Aggregation is hardest because a wrong aggregation can preserve the winning label on ranking tasks. Error bars are task-clustered 95% bootstrap intervals.

163 resamples for confidence intervals. The mutation confidence interval also clusters by base task to
 164 avoid treating multiple mutants of one task as independent.

165 7 Results

166 7.1 Controlled semantic faults

167 Table 1 shows that execution detects none of the semantic mutants, while universal row/column ro-
 168 bustness checks kill only 9. Operator-conditioned relations account for the gain: full MORPH-DA
 169 kills 364 mutants ($\chi^2 = 353$, $n_{01} = 355$, $n_{10} = 0$, $p = 9.4 \times 10^{-79}$). Filter and aggregation
 170 relations produce 346 kills; grouping, hardcoding, time, and statistical relations add 18. Thus the
 171 approximately $40\times$ ratio over universal checks primarily quantifies how weak non-semantic robust-
 172 ness tests are for operator faults.

173 Detection varies by family (Figure 2): hardcoding 85.3%, grouping 81.0%, ranking 76.1%, filtering
 174 67.6%, and aggregation 28.2%. Low aggregation coverage is partly unidentifiability: if both sum
 175 and mean select the same top group, no output-only relation can distinguish them on that instance.

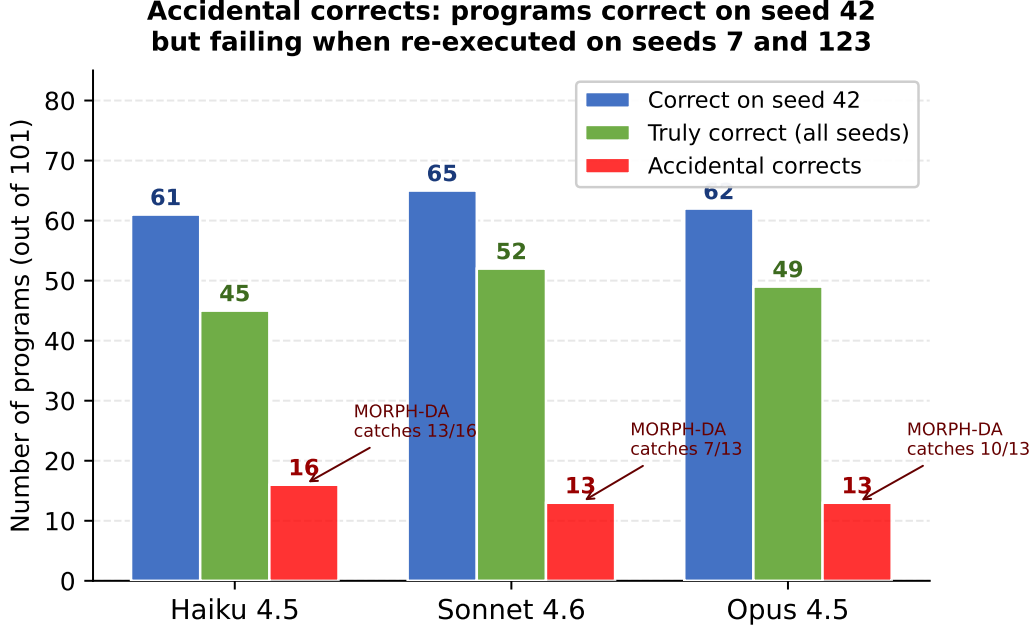


Figure 3: Programs correct on seed 42 split into truly correct and accidental correct after fixed-program re-execution on seeds 7 and 123.

Table 2: Natural-program verification. Condition A uses 101 tasks and naive seed-42 labels. Conditions B/C use the 98-task filter-discriminating set and cross-seed corrected labels. Original execution failures (Haiku 2, Sonnet 0, Opus 2) are excluded from precision/recall/FPR denominators. CIs shown for the primary Condition B.

Model	Condition	Precision	Recall	FPR	F1	AAR
Haiku 4.5	A: naive	58.8	79.0	34.4	67.4	16.7
	B: cross-seed	87.5 [77.1,95.8]	79.2 [67.9,90.6]	14.0 [4.7,25.6]	83.2	22.9
	C: two transformations	86.0	81.1	16.3	83.5	21.7
Sonnet 4.6	A: naive	66.7	72.2	20.0	69.3	16.1
	B: cross-seed	89.2 [78.4,97.3]	67.3 [53.1,79.6]	8.2 [2.0,18.4]	76.7	26.2
	C: two transformations	86.8	67.3	10.2	75.9	26.7
Opus 4.5	A: naive	58.8	81.1	33.9	68.2	14.6
	B: cross-seed	79.2 [68.8,89.6]	79.2 [66.7,89.6]	20.8 [10.4,33.3]	79.2	20.8
	C: two transformations	78.0	81.2	22.9	79.6	19.6

7.2 Accidental correctness

On seed 42, 61/101 Haiku, 65/101 Sonnet, and 62/101 Opus programs appear correct. Re-executing those same programs unchanged on held-out seeds reduces structurally correct counts to 45, 52, and 49 (Figure 3). Therefore, 16, 13, and 13 programs – 26.2%, 20.0%, and 21.0% of apparent successes – are accidental corrects. Without seeing held-out seeds, single-seed MORPH-DA catches 13/16 (81%), 7/13 (54%), and 10/13 (77%) of them.

7.3 Natural-program verification

Under naive labels, correctly flagged accidental programs are counted as false positives, depressing apparent precision to 59–67%. With fixed-program cross-seed ground truth, Condition B obtains 79–89% precision and 67–79% recall (Table 2). Against universal-only relations, the full verifier adds 37–48 true detections and loses none (Holm-corrected $p \leq 3.25 \times 10^{-9}$ for all models).

Table 3: Paired one-step repair on 91 wrong-but-executable Haiku programs from seeds 7, 42, and 123.

ID	Feedback	Fixed	Rate
R0	No retry	0/91	0.0%
R2	Generic “has a bug”	5/91	5.5%
R6	Violated relation name	11/91	12.1%
R7	Full counterexample witness	11/91	12.1%

False positives remain nontrivial: 8.2–20.8%, driven mainly by extreme filter injections, tie-sensitive forced-group tests, and legal computed constants flagged by the hardcoding relation.

A second transformation seed changes recall by +1.9, 0.0, and +2.1 percentage points for Haiku, Sonnet, and Opus, while increasing FPR by about 2 points for each model. It slightly improves F1 for Haiku and Opus but lowers it for Sonnet. Condition C is therefore a multi-transformation sensitivity analysis rather than a default; deployments that value recall over false alarms may choose it.

7.4 One-step repair

The 91-program sample spans seeds 7 (25 programs), 42 (38), and 123 (28), covering 55 distinct tasks. Relation-name and witness feedback each produce 11/91 fixes versus 5/91 for generic feedback (Table 3). For each comparison with R2, $n_{01} = 6$ and $n_{10} = 0$ ($p = 0.041$ unadjusted; $p = 0.082$ after Holm correction for two co-primary comparisons), so we report a higher observed rate rather than a superiority claim. R6 and R7 have the same aggregate rate but repair different programs ($n_{01} = 4$, $n_{10} = 4$, $p = 0.29$); the study does not establish that a richer one-step witness is better than the relation name alone.

Runtime cost. Verification adds no model calls. In a 20-task reference-program pilot measured with raw Python timer timestamps, the full relation set required approximately 40 Python executions per task and had 500 ms mean, 525 ms median, and 1.14 s p95 runtime. This is a small pilot, not a production latency claim.

8 Limitations and Discussion

Specification dependence. MORPH-DA assumes a structured analytical specification. Such specifications may be authored by analysts, derived from governed metric definitions, or extracted by a separate model; automatic extraction and verification of the specification itself are outside scope. The method is gold-answer-free, not supervision-free.

Incomplete necessary conditions. Passing the current relations does not prove correctness. Aggregation coverage is especially limited on label-output tasks when wrong operators preserve the winner. New relation families, richer output contracts, and combinations with learned judges may improve recall.

Benchmark and model scope. The evaluated data are synthetic, all tasks are single-table, and all natural programs come from three variants of one model family using a LangGraph harness. Join and weighted-mean fields are represented in the DSL, but neither capability has implemented relations or evaluated tasks. Generalization to open-weight models, SQL, multi-table analysis, and real enterprise datasets remains to be tested.

Reference and repair limitations. Reference programs and relations share the same specifications, so reference pass rates measure internal consistency rather than an independent FPR. Natural-agent programs provide the independent FPR estimate. Finally, repair success is low and one-step only; equal aggregate R6/R7 rates and non-significant paired differences motivate multi-round and fault-specific repair studies.

225 **Implications for agent verification.** The main lesson is not that metamorphic testing replaces
 226 gold evaluation. Instead, it supplies a reusable environment-grounded signal where direct answers
 227 are unavailable, exposes single-instance accidental correctness, and creates counterexamples that
 228 can support debugging, model-based judging, or human review.

229 9 Conclusion

230 Execution success is an inadequate verifier for analytical agents, and single-instance correctness
 231 can conceal structurally wrong programs. MORPH-DA combines a mutation-grounded benchmark
 232 with specification-conditioned data transformations to test necessary behavioral properties without
 233 revealing the gold output. It detects 64.7% of 563 semantic mutants, identifies a substantial fraction
 234 of accidental corrects, and yields useful precision/recall tradeoffs on natural agent programs without
 235 an additional LLM verification call. The remaining false positives and undetected aggregation errors
 236 reinforce the intended interpretation: metamorphic execution is a complementary runtime verifier,
 237 not a correctness certificate.

238 References

- 239 Steven Cho, Stefano Ruberto, and Valerio Terragni. LLMORPH: Automated metamorphic testing
 240 of large language models. In *Proceedings of the 40th IEEE/ACM International Conference on*
 241 *Automated Software Engineering*, 2025. doi: 10.1109/ASE63991.2025.00385.
- 242 Xueyu Hu, Ziyu Zhao, Shuang Wei, Ziwei Chai, Qianli Ma, Guoyin Wang, Xuwu Wang, Jing Su,
 243 Jingjing Xu, Ming Zhu, Yao Cheng, Jianbo Yuan, Jiwei Li, Kun Kuang, Yang Yang, Hongxia
 244 Yang, and Fei Wu. InfiAgent-DABench: Evaluating agents on data analysis tasks. *arXiv preprint*
 245 *arXiv:2401.05507*, 2024.
- 246 Yue Jia and Mark Harman. An analysis and survey of the development of mutation testing. *IEEE*
 247 *Transactions on Software Engineering*, 37(5):649–678, 2011. doi: 10.1109/TSE.2010.62.
- 248 Yuhang Lai, Chengxi Li, Yiming Wang, Tianyi Zhang, Ruiqi Zhong, Luke Zettlemoyer, Scott Wen-
 249 tau Yih, Daniel Fried, Sida Wang, and Tao Yu. DS-1000: A natural and reliable benchmark for
 250 data science code generation. *arXiv preprint arXiv:2211.11501*, 2022.
- 251 Boyan Li, Zhuowen Liang, Yupeng Xie, Xiaotian Lin, Tianqi Luo, Xinyu Liu, Yizhang Zhu,
 252 Zhangyang Peng, Yuan Li, Zhengxuan Zhang, Jiayi Zhang, Nan Tang, Guoliang Li, and
 253 Yuyu Luo. DataSpace: Benchmarking data agents for verifiable analytics over heterogeneous
 254 workspaces. *arXiv preprint arXiv:2608.03451*, 2026.
- 255 Pingchuan Ma and Shuai Wang. MT-Teql: Evaluating and augmenting consistency of text-to-SQL
 256 models with metamorphic testing. *arXiv preprint arXiv:2012.11163*, 2020.
- 257 Jaehyun Nam, Jinsung Yoon, Jiefeng Chen, Raj Sinha, Jinwoo Shin, and Tomas Pfister. DS-
 258 STAR: Data science agent for solving diverse tasks across heterogeneous formats and open-ended
 259 queries. *arXiv preprint arXiv:2509.21825*, 2026.
- 260 Shuyin Ouyang, Dong Huang, Jingwen Guo, Zeyu Sun, Qihao Zhu, and Jie M. Zhang.
 261 DSCodeBench: A realistic benchmark for data science code generation. *arXiv preprint*
 262 *arXiv:2505.15621*, 2025.
- 263 Goran Petrović, Marko Ivanković, Gordon Fraser, and René Just. Practical mutation testing at scale.
 264 *arXiv preprint arXiv:2102.11378*, 2021.
- 265 Mizanur Rahman, Mohammed Saidul Islam, Ridwan Mahbub, Md Tahmid Rahman Laskar, Shafiq
 266 Joty, and Enamul Hoque Prince. DSAgentBench: Can agents automate end-to-end data-science
 267 workflows in real computer environments? *arXiv preprint arXiv:2608.10366*, 2026.
- 268 Sergio Segura, Gordon Fraser, Ana B. Sanchez, and Antonio Ruiz-Cortes. A survey on metamorphic
 269 testing. *IEEE Transactions on Software Engineering*, 42(9):805–824, 2016. doi: 10.1109/TSE.
 270 2016.2532875.

- 271 Zhaoyan Sun, Shan Zhong, Daizhou Wen, Jiaying Han, Guoliang Li, Ying Yan, Peng Zhang, Yu Su,
272 Xiang Qi, Baolin Sun, Chengyuan Yang, Tao Fang, and Huaiyu Ruan. AgenticDataBench: A
273 comprehensive benchmark for data agents. *arXiv preprint arXiv:2607.01647*, 2026.
- 274 Bo Yang, Yinfen Xia, Weisong Sun, and Yang Liu. Hallucination detection for LLM-based text-to-
275 SQL generation via two-stage metamorphic testing. *arXiv preprint arXiv:2512.22250*, 2025.
- 276 Dan Zhang, Sining Zhoubian, Min Cai, Fengzu Li, Lekang Yang, Wei Wang, Tianjiao Dong, Ziniu
277 Hu, Jie Tang, and Yisong Yue. DataSciBench: An LLM agent benchmark for data science. *arXiv*
278 *preprint arXiv:2502.13897*, 2025.

279 A Additional Experimental Detail

280 A.1 Condition-B confusion matrices

Table 4: Cross-seed corrected confusion matrices. Original execution failures are excluded.

Model	TP	FP	TN	FN	Evaluated
Haiku 4.5	42	6	37	11	96
Sonnet 4.6	33	4	45	16	98
Opus 4.5	38	10	38	10	96

281 A.2 Seed-42 execution breakdown

Table 5: Original seed-42 program outcomes on the 98-task filter-discriminating set.

Model	Correct	Wrong executable	Execution failure
Haiku 4.5	59		37
Sonnet 4.6	62		36
Opus 4.5	61		35

282 A.3 Relation-level false positives

Table 6: Condition-B false positives by dominant relation.

Relation	Haiku	Sonnet	Opus	Typical cause
MR-F1	4	2	6	Imprecise filter logic under extreme rows
MR-G3	1	1	2	Tie-break sensitivity
MR-H1	1	1	2	Legal computed constants

283 A.4 Mutant-generation funnel

284 The verified funnel is 3,030 $\rightarrow$ 709 $\rightarrow$ 563: generated candidates, executable candidates entering
285 equivalence evaluation, and retained non-equivalent mutants, respectively. Of the 709 executable
286 candidates, 146 (20.6%) were equivalent across five oracle seeds and removed. An intermediate
287 `rulemut_stats.json` file from an earlier run is not used for final results.

288 A.5 Repair pairing and provenance

289 The repair log contains 364 rows: 91 source programs crossed with four strategies. Each row in-
290 cludes a stable program identifier, model, data seed, and source hash, enabling row-paired compar-
291 isons. The source-program distribution is 25 programs from seed 7, 38 from seed 42, and 28 from
292 seed 123.

293 A.6 Reproducibility and anonymity

294 The anonymous supplementary package contains task specifications, data generators, relation im-
295 plementations, evaluation scripts, aggregate result tables, and de-identified logs. Public author-
296 identifying repository links are intentionally omitted during double-blind review.